

0304

February 27, 2021

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm.notebook import tqdm

import warnings
warnings.filterwarnings("ignore")

SEED = 123
np.random.seed(SEED)
```

```
[2]: DATA_RAW_PATH = '../data/raw/'
DATA_INTER_PATH = '../data/interim/'
FIGURES = '../figures/'
MODELS = '../models/'
DATA_RAW_NAME = 'teste_smarkio_lbs.xls'
DATA_INTER_NAME = 'df_1.csv'
```

```
[3]: df = pd.read_csv(DATA_INTER_PATH+DATA_INTER_NAME)
df.head(7)
```

```
[3]:
```

	Pred_class	probabilidade	status	True_class
0	2	0.079892	approved	0.0
1	2	0.379377	approved	74.0
2	2	0.379377	approved	74.0
3	2	0.420930	approved	74.0
4	2	0.607437	approved	2.0
5	2	0.690894	approved	2.0
6	2	0.759493	approved	2.0

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 643 entries, 0 to 642
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
```

```

0   Pred_class      643 non-null    int64
1   probabilidade   643 non-null    float64
2   status          643 non-null    object
3   True_class      643 non-null    float64
dtypes: float64(2), int64(1), object(1)
memory usage: 20.2+ KB

```

0.0.1 Questão 3:

Crie um classificador que tenha como output se os dados com status igual a revision estão corretos ou não (Sugestão : Técnica de cross-validation K-fold)

```

[5]: from sklearn.linear_model import LogisticRegression
     from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier,
     ↪ExtraTreesClassifier
     from sklearn.tree import DecisionTreeClassifier
     from sklearn.svm import SVC
     import lightgbm as lgbm
     import xgboost as xgb
     import catboost as ctb

     from sklearn.model_selection import cross_validate, cross_val_score, KFold,
     ↪learning_curve
     from sklearn.metrics import f1_score, recall_score, precision_score,
     ↪classification_report

     import pickle

```

- input: Pred_class, probabilidade
- output: True_class

```

[6]: train = df[df['status'] == 'approved'].drop('status', axis=1)
     test = df[df['status'] == 'revision'].drop('status', axis=1)

     X_train = train.iloc[:,2].values
     y_train = train.iloc[:,2].values

     X_test = test.iloc[:,2].values
     y_test = test.iloc[:,2].values

     print('Treino:',X_train.shape, y_train.shape)
     print('Teste:',X_test.shape, y_test.shape)

```

```

Treino: (600, 2) (600,)
Teste: (43, 2) (43,)

```

```

[7]: models = [
     ('DecisionTree', DecisionTreeClassifier(random_state=SEED)),
     ('RandomForest', RandomForestClassifier(random_state=SEED)),

```

```

('ExtraTree', ExtraTreesClassifier(random_state=SEED)),
('Adaboost', AdaBoostClassifier(random_state=SEED)),
('XGBoost', xgb.XGBClassifier(random_state=SEED)),
('LightGBM', lgbm.LGBMClassifier(random_state=SEED)),
('Catboost', ctb.CatBoostClassifier(random_state=SEED, verbose=False)),
('LogisticRegression', LogisticRegression(random_state=SEED)),
('SVC', SVC(random_state=SEED))
]

```

```

[8]: original = pd.DataFrame()

for name, model in tqdm(models):

    kfold = KFold(n_splits=3, random_state=SEED, shuffle=True)
    score = cross_validate(model, X_train, y_train, cv=kfold,
→scoring=['precision_weighted', 'recall_weighted', 'f1_weighted'],
→return_train_score=True)
    additional = pd.DataFrame({
        'precision_train': np.mean(score['train_precision_weighted']),
        'precision_test' : np.mean(score['test_precision_weighted']),
        'recall_train': np.mean(score['train_recall_weighted']),
        'recall_test' : np.mean(score['test_recall_weighted']),
        'f1_train' : np.mean(score['train_f1_weighted']),
        'f1_test' : np.mean(score['test_f1_weighted']),
    }, index=[name])

    new = pd.concat([original, additional], axis=0)
    original = new

```

```
0%|          | 0/9 [00:00<?, ?it/s]
```

```

[18:03:17] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.3.0/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default
evaluation metric used with the objective 'multi:softprob' was changed from
'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the
old behavior.

```

```

[18:03:18] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.3.0/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default
evaluation metric used with the objective 'multi:softprob' was changed from
'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the
old behavior.

```

```

[18:03:19] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.3.0/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default
evaluation metric used with the objective 'multi:softprob' was changed from
'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the
old behavior.

```

0.0.2 Questão 4:

Compare três métricas de avaliação aplicadas ao modelo e descreva sobre a diferença.

```
[9]: original
```

```
[9]:
```

	precision_train	precision_test	recall_train \
DecisionTree	0.997536	0.593028	0.997500
RandomForest	0.997557	0.539982	0.997500
ExtraTree	0.997536	0.546556	0.997500
Adaboost	0.110871	0.118734	0.222500
XGBoost	0.982366	0.619289	0.988333
LightGBM	0.984414	0.590629	0.983333
Catboost	0.973669	0.572989	0.973333
LogisticRegression	0.105122	0.055064	0.215000
SVC	0.104726	0.092244	0.270833

	recall_test	f1_train	f1_test
DecisionTree	0.593333	0.997498	0.573649
RandomForest	0.553333	0.997486	0.525923
ExtraTree	0.561667	0.997498	0.533728
Adaboost	0.225000	0.124007	0.128146
XGBoost	0.633333	0.984853	0.605545
LightGBM	0.610000	0.982171	0.580296
Catboost	0.593333	0.972207	0.559784
LogisticRegression	0.185000	0.105144	0.075540
SVC	0.241667	0.148538	0.129543

Como escolher uma dessas métricas?

Dependeria do problema em si do negócio. - Caso os falsos positivos impactassem mais no negócio a escolha seria: *precision*. - Caso os falsos negativos tivessem maior impacto: *recall*. - Caso em que o negócio precisa de um equilíbrio entre essas duas métricas citadas acima: *f1-score*.

Como existe um equilíbrio entre *Precision* e *Recall* nesse caso, irei avaliar *f1-Score*.

Podemos perceber um possível problema com *overfitting*. O modelo está se saindo muito bem no treino, porém no teste não.

```
[10]: model = xgb.XGBClassifier(random_state=SEED)
```

```
[11]: train_sizes, train_scores, test_scores= learning_curve(model, X_train, y_train,
    ↪cv=kfold, n_jobs=-1, train_sizes=np.linspace(.1, 1.0, 5))

train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

fig, ax = plt.subplots(figsize=(15, 10))
```

```


ax.grid()
ax.fill_between(train_sizes, train_scores_mean - train_scores_std,
    ↪ train_scores_mean + train_scores_std, alpha=0.1, color="r")
ax.fill_between(train_sizes, test_scores_mean - test_scores_std,
    ↪ test_scores_mean + test_scores_std, alpha=0.1, color="g")

ax.plot(train_sizes, train_scores_mean, 'o-', color="r", label="Training score")
ax.plot(train_sizes, test_scores_mean, 'o-', color="g", label="Cross-validation
    ↪ score")

ax.legend(loc="best")
ax.set_xlabel("Training examples")
ax.set_ylabel("Score")

fig.savefig(FIGURES+'learning_curve.png')

```



Enfrentamos um problema de *overfitting*

```

[12]: results = []
names = []

for name, model in tqdm(models):

```

```

    cv_results = cross_val_score(model, X_train, y_train, cv=kfold,
    ↪scoring=('f1_weighted'))
    results.append(cv_results)
    names.append(name)

fig = plt.figure(figsize=(12,6))
fig.suptitle('Comparação entre algoritmos - F1-Score')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.ylabel('Precision score')
plt.xticks(rotation=45)
plt.grid(b=False)
plt.show()

fig.savefig(FIGURES+'/comparacao-modelos.png')

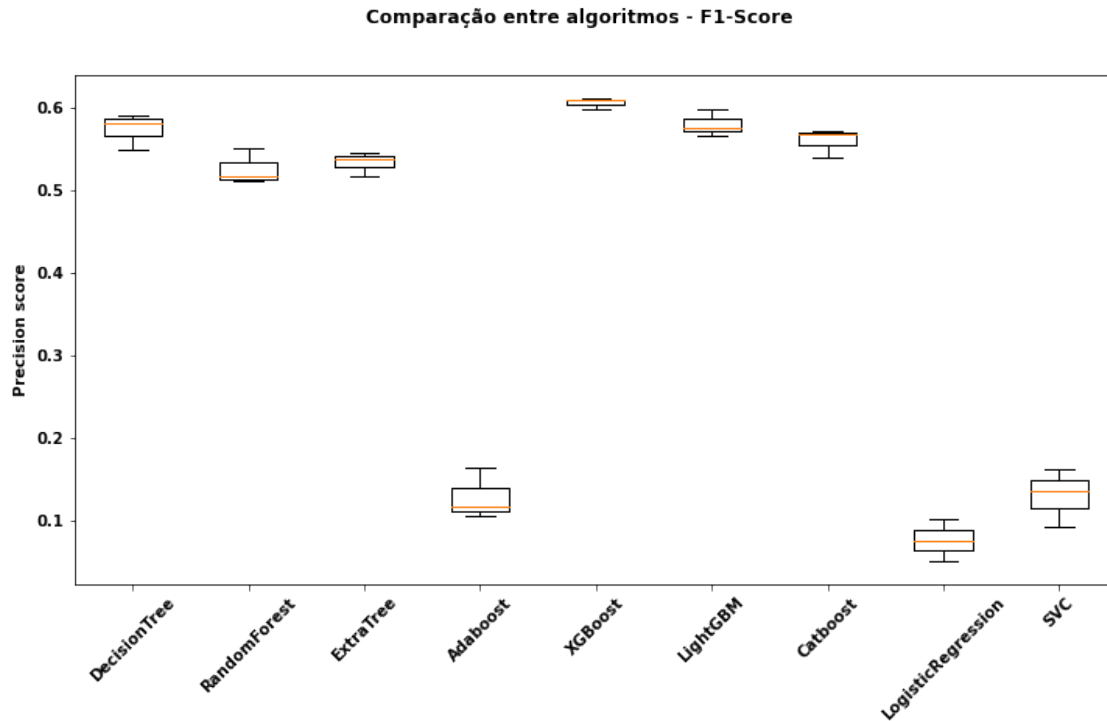
```

0%| | 0/9 [00:00<?, ?it/s]

[18:06:29] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.0/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[18:06:31] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.0/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[18:06:32] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.0/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.



Como existem várias *classes* na previsão, utilizo todas as métricas com `*'_weighted'`, *assim consigo, com base na proporção da classe** que está sendo prevista, colocar um ponderar/colocar um peso sobre essa *classe*.

Modelo selecionado: *XGBoost*

Predict

```
[13]: model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	0
2.0	0.00	0.00	0.00	3
3.0	0.27	1.00	0.43	3
4.0	0.00	0.00	0.00	3
11.0	0.00	0.00	0.00	1
12.0	0.00	0.00	0.00	1
17.0	0.00	0.00	0.00	1
22.0	0.00	0.00	0.00	1
24.0	0.00	0.00	0.00	5
25.0	0.00	0.00	0.00	3

32.0	0.00	0.00	0.00	1
36.0	0.00	0.00	0.00	1
39.0	0.00	0.00	0.00	2
43.0	0.00	0.00	0.00	2
55.0	0.00	0.00	0.00	2
60.0	0.00	0.00	0.00	4
74.0	0.00	0.00	0.00	0
77.0	0.00	0.00	0.00	3
81.0	0.00	0.00	0.00	1
84.0	0.00	0.00	0.00	1
86.0	0.00	0.00	0.00	1
96.0	1.00	1.00	1.00	2
110.0	0.00	0.00	0.00	0
113.0	0.00	0.00	0.00	1
114.0	0.00	0.00	0.00	1
accuracy			0.12	43
macro avg	0.05	0.08	0.06	43
weighted avg	0.07	0.12	0.08	43

Modelo conseguiu acertar 56% dos dados que estavam em *revision*.

```
[14]: pickle.dump(model, open(MODELS+'modelo_classificador.sav', 'wb'))
```

Exportação do modelo.